

CS 228: Introduction to Data Structures

Lecture 39

Shortest Paths Trees

Throughout this lecture, G denotes a weighted graph — directed or undirected —, s denotes the source node in G , and for each edge (u,v) , $w(u,v)$ is the weight (length) of e . For each node v in G , $\text{dist}(v)$ denotes the length of a shortest path from s to v in G . Note that there may be multiple shortest paths from s to v , but they all have the same total length $\text{dist}(v)$. Also, since edge lengths are positive, $\text{dist}(s) = 0$.

A ***shortest-path tree*** is a tree T made up nodes and edges of G and rooted at s , such that the path from s to any other vertex v in T is a shortest path from s to v in G . For instance, a BFS tree is a shortest-path tree for the case where all edges have the same length.

To see that a shortest-path tree always exists, suppose that we know $\text{dist}(v)$ for every node v in G reachable from s . Let P be any shortest paths from s to v . Let $\text{pred}(v)$ be the last node in P before v . Let P' be the subpath of P that goes from s to $\text{pred}(v)$. Then, P' must be a shortest path from s to $\text{pred}(v)$ — if it were not, P would not be a

shortest path to v , violating our assumption. This means that

$$\text{dist}(v) = \text{dist}(\text{pred}(v)) + w(\text{pred}(v), v).$$

Since each node other than s has exactly one predecessor, the set of edges $(\text{pred}(v), v)$, $v \neq s$, form a tree rooted at s . By construction, if the tree contains a path from s to v , this path must be a shortest path from s to v in G .

The length of the shortest path from s to any other node v is unique; the path itself, however, may not be. That is, there may be multiple shortest paths (all of the same length) from s to v . Draw some examples to verify this!

Dijkstra's algorithm

Dijkstra's algorithm starts from the source node s , and in each iteration adds another node to the shortest-path tree. This node is the point closest to the root that is still outside the tree. The process resembles breadth-first search using a priority queue instead of a FIFO queue. Note, however, that it is *not* breadth-first search: we do not care about the number of edges on the tree path, only the sum of their weights.

Throughout its execution, Dijkstra's algorithm maintains the length $d(v)$ of the best path from s to v found so far,

for each node v in G . At all times, $d(v) \geq \text{dist}(v)$ (where, as before, $\text{dist}(v)$ is the length of the shortest path from s to v). Initially

- $d(s) = 0$, and
- $d(v) = \infty$ for every node v different from s .

Dijkstra's algorithm also keeps track of v 's predecessor $\text{pred}(v)$ in the current best path from s to v . Initially, $\text{pred}(v)$ is null for every node v .

Dijkstra's algorithm maintains a set C of nodes, which we will call the **cloud**. The algorithm repeats the following two steps until all nodes are in the cloud.

Step 1. Pick a node u such that $d(u)$ is minimum among the nodes not in C and add u to C .

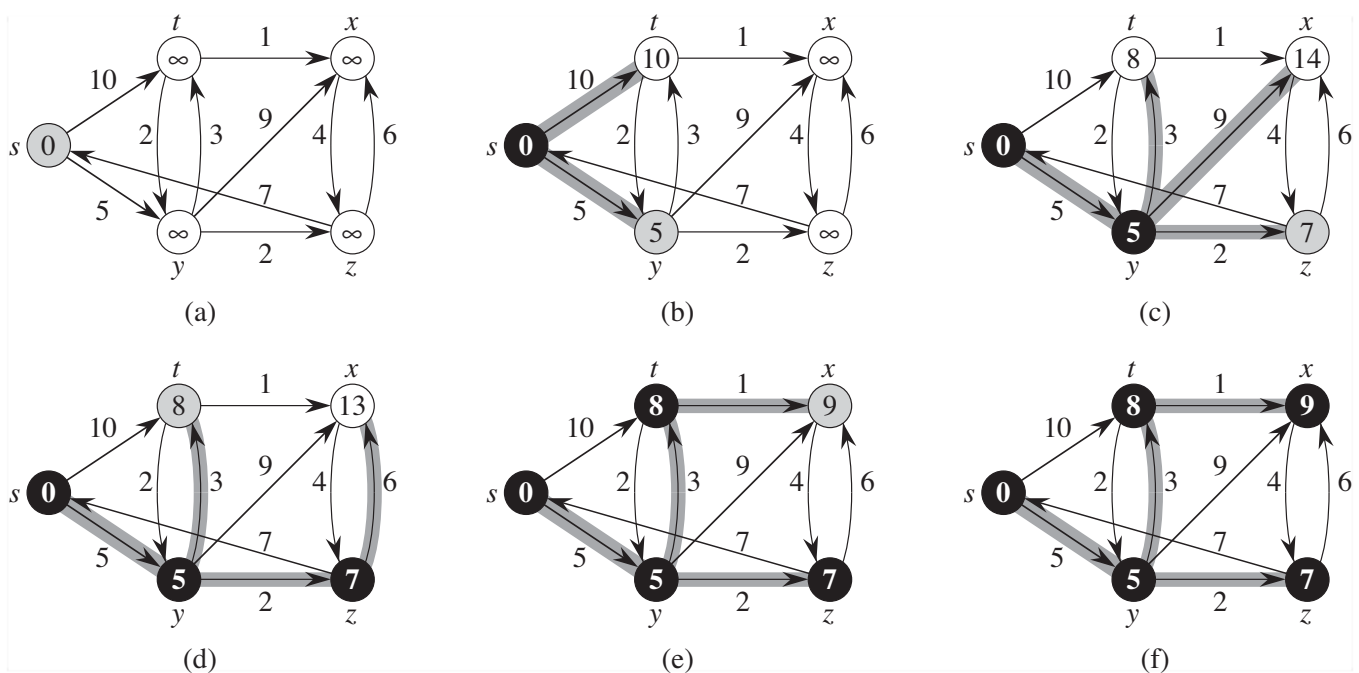
Step 2. For each neighbor v of u , **relax** edge (u,v) . That is, if

$$d(v) > d(u) + w(u,v), \quad (\dagger)$$

make $\text{pred}(v) = u$ and $d(v) = d(u) + w(u,v)$.

The intuitive justification of Step 2 is that encountering condition (\dagger) implies that we have just found a path of length less than $d(v)$ to v via u , so we should update $d(v)$.

Example¹. In the sample execution of Dijkstra's algorithm shown below black nodes are in the cloud. The node u with the minimum d -value at each step is shaded. Shaded edges indicate predecessor values.



Correctness

The correctness of Dijkstra's algorithm is a consequence of the following lemma. As before, $\text{dist}(v)$ denotes the length of the shortest path from s to v .

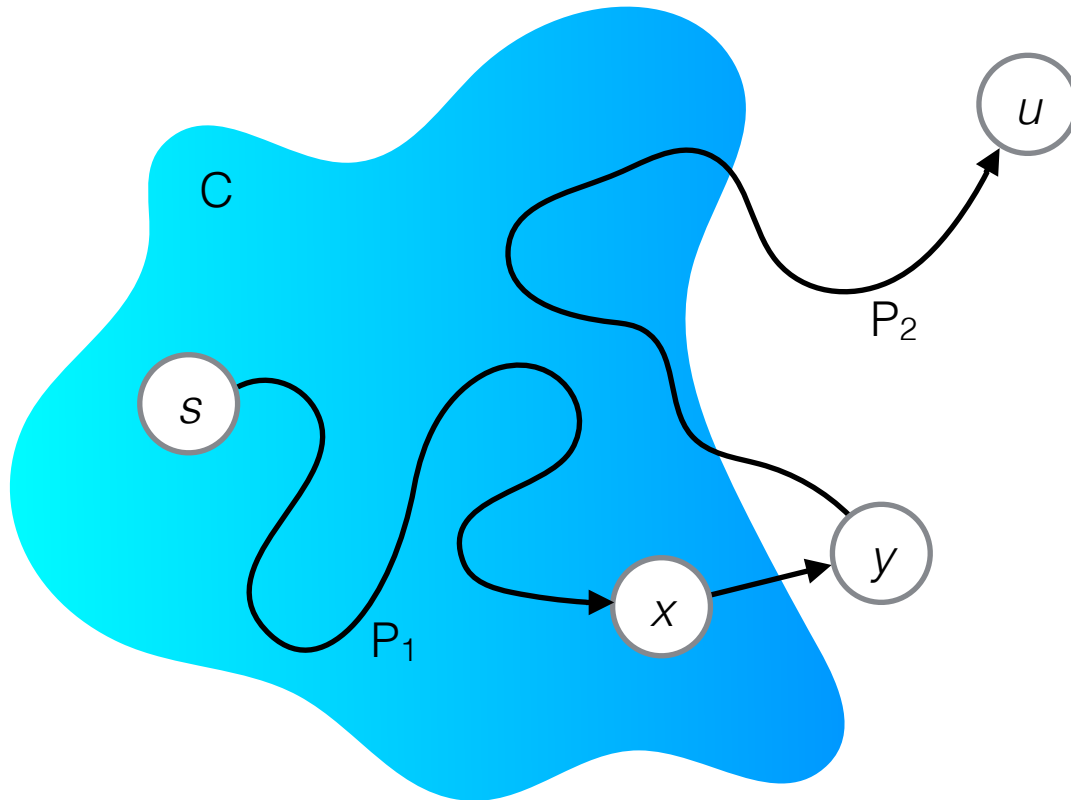
¹ This example is taken from CLRS. Slides for this example and another example, which uses an undirected graph, are posted on Blackboard.

Lemma. *At the start of each iteration of the algorithm, $d(v) = \text{dist}(v)$ for every node v in C .*

Proof Sketch. The lemma is trivially true at the start, because C is empty. The first iteration adds s to C . Since $\text{dist}(s) = d(s) = 0$, the lemma remains true. (You can also check that it's true after the second iteration.) Let's assume that a number of iterations have taken place and that, so far, $d(v) = \text{dist}(v)$ for every node v in C . Now, suppose that in the next iteration, the algorithm makes a mistake, and adds to C a node u such that $\text{dist}(u) \neq d(u)$. We'll sketch why this mistake cannot happen.

Suppose P is a shortest path from s to u . Let y be the first node in P that is not in C and let x be the node immediately preceding y in P . Then, as the figure on the next page illustrates, we can decompose P into three parts: (1) a path P_1 from s to x , (2) the edge (x,y) , and (3) a path P_2 from y to u . The key observation is that we would then have $d(y) < d(u)$ and so node u would not have been chosen by Dijkstra's algorithm, a contradiction.

We emphasize that the preceding argument is just a sketch of a proof. For a formal proof, consult CLRS.



Implementation

The pseudocode for Dijkstra's algorithm shown on the next page uses a priority queue Q , implemented as a binary heap, to store the nodes of G that are not in the cloud. Nodes in Q are prioritized by d -values, so the `removeMin()` operation executed at the beginning of each iteration always extracts the node in the cloud with the minimum d -value.

Relaxing edge (u,v) may necessitate decreasing $d(v)$, so Q requires a `decreaseKey()` method (specifically, we need this in **(*)**). This method has two steps:

1. Reduce the value of $d(v)$ in Q .
2. Use `percolateUp()` to readjust Q .

The time for this is dominated by `percolateUp()`, which is logarithmic in the number of nodes in the heap; i.e., it is $O(\log V)$. We leave the details as an exercise.

```
DIJKSTRA( $G, s$ ):  
    for each node  $v$  in  $G$  except  $s$   
         $d(v) = \infty$   
         $\text{pred}(v) = \text{null}$   
     $d(s) = 0$   
    let  $Q$  be a priority queue containing  
        the nodes of  $G$  prioritized by  $d$ -value  
    let  $C$  be an empty set  
    while  $Q$  is not empty  
         $u = \text{removeMin}(Q)$   
         $C = C \cup \{u\}$   
        foreach neighbor  $v$  of  $u$   
            // Relax ( $u, v$ )  
            if  $d(v) > d(u) + w(u, v)$   
                 $d(v) = d(u) + w(u, v)$  (*)  
                 $\text{pred}(v) = u$   
    return  $\text{pred}$ 
```

Time complexity. Assume that we use a binary heap to implement the priority queue Q . Initializing Q and the d - and pred -values takes $O(V)$ time. The **while** loop does the

bulk of the work. There are V iterations. Removing the minimum element at the beginning of each iteration takes $O(\log V)$ time, for a total of $O(V \log V)$ time over all V iterations. During these iterations, the number of times we may have to update $\text{pred}(v)$ and reduce the value of $d(v)$ in Q is at most equal to the in-degree of v (assume that the graph is directed — a similar argument applies to undirected graphs). The total number of such updates over all iterations is at most the sum of the in-degrees; this sum equals E . Each update takes $O(\log V)$ time (dominated by the time to update Q), so the total time spent on these steps is $O(E \log V)$. This is the dominant term in the run time, so the time complexity of Dijkstra's algorithm is $O(E \log V)$.

Single-Source Shortest Paths in DAGs

In several important applications of the single-source shortest path problem, the graph G is a weighted directed acyclic graph (DAG). It turns out that for DAGs there is a faster and simpler alternative to Dijkstra's algorithm, which uses edge relaxation (like Dijkstra's algorithm) and topological sorting. This algorithm achieves a running time of $O(V + E)$, without using any major data structures — in particular, no heaps. Before describing the DAG algorithm, we need to review a few things about shortest paths that apply to arbitrary graphs, not just acyclic ones.

Recall that by **relaxing** an edge (u,v) in a graph G we mean (1) checking whether there is a path to v going via u that is better than the current best path to v and (2) if the answer is “yes”, updating $d(v)$ and $\text{pred}(v)$ accordingly.

```
RELAX( $u, v$ ):  
|   if  $d(v) > d(u) + w(u, v)$   
|       |    $d(v) = d(u) + w(u, v)$   
|       |    $\text{pred}(v) = u$ 
```

Suppose G is a weighted graph (not necessarily a DAG), where edge weights may be negative weights, but which has no negative cycles (i.e., cycles where the sum of the edge weights is negative). Now, suppose we put $d(s) = 0$ and $d(v) = \infty$ for all other nodes v . Then, the following property can be proved.

Path-relaxation property. If $p = v_0, v_1, \dots, v_k$ is a shortest path from $s = v_0$ to v_k , and we relax the edges of p in the order $(v_0, v_1), (v_1, v_2), \dots, (v_{k-1}, v_k)$, then $d(v_k) = \text{dist}(v_k)$. This property holds regardless of any other relaxation steps that occur, even if they are intermixed with relaxations of the edges of p .

Now suppose G is a DAG, and that we have sorted G topologically. The shortest path algorithm shown next begins by initializing the d - and pred -values as usual. It then makes a single pass over the nodes in the

topologically order. As it process each node, it relaxes each edge that leaves the node.

```
DAGSHORTESTPATHS( $G, s$ ):
```

```
    topologically sort the nodes of  $G$ 
```

```
    // Initialization
```

```
    foreach node  $v$  in  $G$ 
```

```
         $d(v) = \infty$ 
```

```
         $\text{pred}(v) = \text{null}$ 
```

```
     $d(s) = 0$ 
```

```
    foreach node  $u$  in  $G$  in topological order
```

```
        foreach neighbor  $v$  of  $u$ 
```

```
            // Relax ( $u, v$ )
```

```
            if  $d(v) > d(u) + w(u, v)$ 
```

```
                 $d(v) = d(u) + w(u, v)$ 
```

```
                 $\text{pred}(v) = u$ 
```

```
    return  $\text{pred}$ 
```


We claim that at termination, for every node v , $d(v)$ is the length of the shortest path from s to v . To verify this, observe that if G contains a path from node u to node v , then u precedes in the topological sort. Thus, for every node v , the edges of a shortest path from s to v are relaxed in order, from first to last. By the path-relaxation

property, this means that, at termination, $d(v)$ is indeed the length of a shortest path from s to v .

Observe that DAGSHORTESTPATHS is correct even for DAGs with negative-weight edges. The reason is that shortest paths are always well-defined in a DAG, since even if there are negative-weight edges, no negative-weight cycles can exist. Further, the path relaxation property holds as long as there are no negative cycles.

Example². The figure on the next page shows a sample execution of DAGSHORTESTPATHS. The nodes are topologically sorted from left to right. The source node is s . The d values appear within the nodes, and shaded edges indicate the pred-values. The newly blackened node in each iteration was used as u in that iteration.

² This example is from CLRS and slides for it are posted on Blackboard.



Time complexity. As we have seen, topological sort takes $O(V + E)$ time. The initialization of d- and pred-values takes $O(V)$ time. The outer **foreach** loop makes one iteration per vertex. Altogether, the inner **foreach** loop relaxes each edge exactly once. Because each iteration of the inner **foreach** loop takes $O(1)$ time, the total running time is $O(V + E)$, which is linear in the size of an adjacency-list representation of the graph.

Caution: Shortest paths are *not* well-defined when negative cycles are present. In fact, if G has a negative cycle C , you can make the length of a path between two nodes in C arbitrarily small by going around C multiple times.